

PROJECT - CALIFORNIA HOUSING PRICE PREDICTION

DESCRIPTION:

The US Census Bureau has published California Census Data which has 10 types of metrics such as the population, median income, median housing price, and so on for each block group in California. The dataset also serves as an input for project scoping and tries to specify the functional and nonfunctional requirements for it.

BACKGROUND OF THE PROBLEM STATEMENT:

The project aims at building a model of housing prices to predict median house values in California using the provided dataset. This model should learn from the data and be able to predict the median housing price in any district, given all the other metrics. Districts or block groups are the smallest geographical units for which the US Census Bureau Publishes sample data (a block group typically has a population of 600 to 3,000 people). There are 20,640 districts in the project dataset.

DATASET DESCRIPTION: (20640 rows x 10 columns)

FIELD	DESCRIPTION
longitude	(signed numeric - float) : Longitude value for the block in California, USA
latitude	(numeric - float) : Latitude value for the block in California, USA
housing_median_age	(numeric - int) : Median age of the house in the block
total_rooms	(numeric - int) : Count of the total number of rooms (excluding bedrooms) in all houses in the block
total_bedrooms	(numeric - float) : Count of the total number of bedrooms in all houses in the block
population	(numeric - int) : Count of the total number of population in the block
households	(numeric - int) : Count of the total number of households in the block
median_income	(numeric - float) : Median of the total household income of all the houses in the block
ocean_proximity	(numeric - categorical) : Type of the landscape of the block [Unique Values : 'NEAR BAY', '<1H OCEAN', 'INLAND', 'NEAR OCEAN', 'ISLAND']
median_house_value	(numeric - int) : Median of the household prices of all the houses in the block

[PLEASE DOWNLOAD THE DATASET FROM HERE](#) 

QUESTIONS TO BE ANSWERED WITH ANALYSIS:

1. **Build a model** of housing prices to predict median house values in California using the provided dataset.
2. **Train the model** to learn from the data to predict the median housing price in any district, given all the other metrics.
3. **Predict housing prices** based on median_income and plot the regression chart for it.

PROJECT GUIDELINES:

1. LOAD THE DATA:

- 1.1. Read the “**housing.csv**” file from the folder into the program.
- 1.2. Print first few rows of this data.
- 1.3. Extract input (X) and output (Y) data from the dataset.

2. HANDLE MISSING VALUES:

- 2.1. Fill the missing values with the mean of the respective column.

3. ENCODE CATEGORICAL DATA:

- 3.1. Convert categorical column in the dataset to numerical data.

4. SPLIT THE DATASET:

- 4.1. Split the data into 80% training dataset and 20% test dataset.

5. STANDARDIZE DATA:

- 5.1. Standardize training and test datasets.

6. PERFORM LINEAR REGRESSION:

- 6.1. Perform Linear Regression on training data.
- 6.2. Predict output for test dataset using the fitted model.
- 6.3. Print root mean squared error (RMSE) from Linear Regression.

[HINT: Import **mean_squared_error** from **sklearn.metrics**]

7. PERFORM DECISION TREE REGRESSION:

- 7.1. Perform Decision Tree Regression on training data.
- 7.2. Predict output for test dataset using the fitted model.
- 7.3. Print root mean squared error from Decision Tree Regression.

8. PERFORM RANDOM FOREST REGRESSION:

- 8.1. Perform Random Forest Regression on training data.
- 8.2. Predict output for test dataset using the fitted model.
- 8.3. Print RMSE (root mean squared error) from Random Forest Regression.

[BONUS EXERCISE]

9. PERFORM LINEAR REGRESSION WITH ONE INDEPENDENT VARIABLE:

- 9.1. Extract just the median_income column from the independent variables (From **X_train** and **X_test**).
- 9.2. Perform Linear Regression to predict housing values based on **median_income**.
- 9.3. Predict output for test dataset using the fitted model.
- 9.4. Plot the fitted model for training data as well as for test data to check if the fitted model satisfies the test data.